

Inventory Management System for B2B SaaS

Name – Aakash Raju Mohole

Email – aakashmohole@gmail.com

Github – <https://github.com/aakashmohole/Bynry-Case-Study-Task/tree/main>

PART 1 - Code Review & Debugging — Approach & Analysis

1. Issues in the Provided code

Technical Issues –

- No check before creating product; duplicate SKUs could be created.
- Two commits used instead of one atomic transaction, risking database inconsistency if one commit fails.
- Decimal price – The code trust that data['price'] is valid this price should be convert to decimal.
- Implement error handling if something fails it crashes or corrupts all data.
- It is unclear if the Product model is meant to be per-warehouse; if not, the logic here can duplicate products unnecessarily.

Business Logic Issues

- The design suggests a new product per warehouse, but product should have a single entry with inventory per warehouse.
- If product creation or inventory insertion in case fails, IDs might become out of sync and this create problem.
- If the app should support multiple companies, there's no separation implemented we have to implement this.
- If inventory creation fails after product is added, the product remains without inventory—a broken state.

2. Impact of Each Issue in Production

- Duplicate SKUs lead to confusion in inventory tracking, order processing, and reporting and data will become redundant.
- Partial failures create split states, making rollback and consistency difficult.
- Invalid data can crash the server or lead to incomplete/corrupt entries.
- Inaccurate prices (due to float imprecision) can cause billing and accounting errors.
- Any exception causes 500 errors and possible data leaks or losses.
- Duplicating product entries for each warehouse breaks analytics, reporting, and increases maintenance burden.

3. Correct Code Version with explanations

```
from flask import request, jsonify
from sqlalchemy.exc import IntegrityError
from decimal import Decimal

@app.route('/api/products', methods=['POST'])
def create_product():
    data = request.json

    # Validate required fields
    required_fields = ['name', 'sku', 'price', 'warehouse_id']
    for field in required_fields:
        if field not in data:
            return jsonify({"error": f"Missing required field: {field}"}), 400

    # Parse and validate price as Decimal
    try:
        price = Decimal(str(data['price']))
    except:
        return jsonify({"error": "Invalid price format"}), 400

    # Check SKU uniqueness
    if Product.query.filter_by(sku=data['sku']).first():
        return jsonify({"error": "SKU already exists"}), 409

    # Begin transaction
    try:
        product = Product(
            name=data['name'],
```

```

        sku=data['sku'],
        price=price
        # exclude warehouse_id, since product is global across warehouses
    )
    db.session.add(product)
    db.session.flush() # gets the product.id before commit

    inventory = Inventory(
        product_id=product.id,
        warehouse_id=data['warehouse_id'],
        quantity=data['initial_quantity']
    )
    db.session.add(inventory)
    db.session.commit()
except IntegrityError as e:
    db.session.rollback()
    return jsonify({"error": "Database error: " + str(e.orig)}), 500
except Exception as e:
    db.session.rollback()
    return jsonify({"error": str(e)}), 500

return jsonify({"message": "Product created", "product_id": product.id}), 201

```

PART 2: Database Design

1. Database Schema Design

Entity	Attributes	Relationships
Company	id (PK), name, created_at	1-to-many with Warehouse, 1-to-many with Product
Warehouse	id (PK), company_id (FK), name, location, created_at	Many-to-1 with Company, 1-to-many with Inventory
Product	id (PK), company_id (FK), name, sku, price, is_bundle	Many-to-1 with Company, many-to-many with Supplier, many-to-many with Product (as bundle), 1-to-many with Inventory
Supplier	id (PK), name, contact_info, created_at	many-to-many with Product via Product_Supplier
Product_Supplier	product_id (PK, FK), supplier_id (PK, FK)	Many-to-1 with Product, many-to-1 with Supplier
Product_Bundle	bundle_id (PK, FK), product_id (PK, FK), quantity	Many-to-1 with Product (bundle and component product)
Inventory	id (PK), product_id (FK), warehouse_id (FK), quantity	Many-to-1 with Product, many-to-1 with Warehouse, 1-to-many with Inventory_Change

	id (PK), inventory_id (FK), change_amount, reason, changed_at, changed_by	
Inventory_Change		Many-to-1 with Inventory

- **SQL DDL**

-- **1. Companies**

```
CREATE TABLE company (
  id SERIAL PRIMARY KEY,
  name VARCHAR(255) NOT NULL,
  -- Add fields like address, contact info, etc.
  created_at TIMESTAMP NOT NULL DEFAULT NOW()
);
```

-- **2. Warehouses**

```
CREATE TABLE warehouse (
  id SERIAL PRIMARY KEY,
  company_id INTEGER NOT NULL REFERENCES company(id) ON DELETE
  CASCADE,
  name VARCHAR(255) NOT NULL,
  location VARCHAR(255),
  created_at TIMESTAMP NOT NULL DEFAULT NOW()
);
```

-- **3. Products**

```
CREATE TABLE product (
  id SERIAL PRIMARY KEY,
  company_id INTEGER NOT NULL REFERENCES company(id) ON DELETE
  CASCADE,
  name VARCHAR(255) NOT NULL,
  sku VARCHAR(64) NOT NULL,
  price DECIMAL(12, 2) NOT NULL,
  is_bundle BOOLEAN NOT NULL DEFAULT FALSE,
  -- Additional fields: description, etc.
  UNIQUE (company_id, sku) -- SKU unique within company
);
```

-- **4. Bundled Products (For bundles containing other products)**

```
CREATE TABLE product_bundle (  
    bundle_id INTEGER NOT NULL REFERENCES product(id) ON DELETE  
    CASCADE,  
    product_id INTEGER NOT NULL REFERENCES product(id) ON DELETE  
    CASCADE,  
    quantity INTEGER NOT NULL DEFAULT 1,  
    PRIMARY KEY (bundle_id, product_id)  
);
```

-- 5. Suppliers

```
CREATE TABLE supplier (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR(255) NOT NULL,  
    contact_info VARCHAR(255),  
    -- Add address, phone, etc.  
    created_at TIMESTAMP NOT NULL DEFAULT NOW()  
);
```

-- 6. Product - Supplier mapping (Products can have one or more suppliers)

```
CREATE TABLE product_supplier (  
    product_id INTEGER NOT NULL REFERENCES product(id) ON DELETE  
    CASCADE,  
    supplier_id INTEGER NOT NULL REFERENCES supplier(id) ON DELETE  
    CASCADE,  
    PRIMARY KEY (product_id, supplier_id)  
);
```

-- 7. Inventory (quantity per product per warehouse)

```
CREATE TABLE inventory (  
    id SERIAL PRIMARY KEY,  
    product_id INTEGER NOT NULL REFERENCES product(id) ON DELETE  
    CASCADE,  
    warehouse_id INTEGER NOT NULL REFERENCES warehouse(id) ON  
    DELETE CASCADE,  
    quantity INTEGER NOT NULL DEFAULT 0,  
    UNIQUE (product_id, warehouse_id)  
);
```

-- 8. Inventory Changes (Audit/history log)

```
CREATE TABLE inventory_change (
    id SERIAL PRIMARY KEY,
    inventory_id INTEGER NOT NULL REFERENCES inventory(id) ON DELETE
    CASCADE,
    change_amount INTEGER NOT NULL,
    reason VARCHAR(255),
    changed_at TIMESTAMP NOT NULL DEFAULT NOW(),
    changed_by INTEGER -- Optionally FK to user table
);
```

2. Identify Gaps / Product Questions

- Will we implement user access control to define which users can view or manage specific company data?
- How should we model suppliers? Can a single supplier provide products to multiple companies, or should supplier relationships be limited to one company? In other words, are suppliers global entities or company-specific?
- Regarding product bundles, can a bundle include other bundles (creating a nested structure), or are bundles composed solely of individual products?
- How will inventory for bundles be handled? Will bundles have their own distinct inventory, or will their availability be derived from the inventory of their constituent components?
- What types of inventory transactions do we anticipate tracking? (e.g., sales, returns, damages, restocks, transfers).
- What is the desired scope of SKU uniqueness? Should SKUs be unique across the entire system globally, within each company, or per warehouse? Any difference between warehouse types? (online, physical, drop-ship, etc.)
- Do we need to differentiate between various warehouse types (e.g., online fulfilment centres, physical retail stores, drop-ship locations)? If so, what are the key distinctions?
- Should we store specific warehouse attributes such as capacity, physical address, or other relevant details?
- Is it necessary to track product pricing on a per-supplier basis?

3. Explain Decisions:

Indexes and Constraints –

a) Indexes and Constraints:

Unique (company_id, sku):

- Keeps SKUs unique within a company—prevents duplicates but allows same SKU in different companies (fits multi-tenancy).

Foreign Keys:

- Ensure no orphaned warehouses, products, inventory, etc.

Unique (product_id, warehouse_id) for Inventory:

- Guarantees one inventory record per product/location.

Primary Keys/Composite Keys in Mapping Tables:

- Supports fast JOIN operations and avoids duplicates.

b) Audit Table (inventory_change):

- Tracks inventory adjustments over time for history and troubleshooting.

c) Bundles (product_bundle):

- Supports "products as bundles" with flexible contents; easy to navigate for inventory checks or UI displays.

d) Scalability:

- Indexed lookup on SKU, company/product IDs.
- Avoids denormalization; easy to extend (add more product attributes or relationships later).

e) Flexibility:

- Schema supports per-company isolation (key for B2B SaaS/multi-tenant).
- Bundles and product-supplier associations are generic; can accommodate changing requirements without major refactoring.

PART 3: API Implementation

Assumptions & Schema Extensions

To implement the endpoint, I assume:

- There is a table sales with columns:
 - id, product_id, warehouse_id, quantity, sale_date.
- Product types are implied by a product_type field in product table (e.g. "standard", "bundle").
- Low-stock thresholds are stored in a table product_type_threshold (product_type → threshold stock level).
- supplier table has a contact_email column.
- Only sales within the last 30 days count as recent sales activity.
- Days until stockout calculated as:
 - $\text{current_stock} / \text{average_daily_sales}$ for that product across the company (all warehouses).
- If average_daily_sales = 0, set days_until_stockout = null or very high

API CODE -

```
from flask import Flask, jsonify
```

```
from flask_sqlalchemy import SQLAlchemy
```

```
from datetime import datetime, timedelta
```

```
from sqlalchemy import func, Boolean
```

```
from flasgger import Swagger
```

```
from dotenv import load_dotenv
```

```
import os
```

```
load_dotenv()
```

```
app = Flask(__name__)
```

```
swagger = Swagger(app, config=None, template=None)
```

```
app.config['SQLALCHEMY_DATABASE_URI'] = os.getenv('DATABASE_URL')
```

```
app.config['SQLALCHEMY_TRACK_MODIFICATIONS'] = False
```

```
print("SQLALCHEMY_DATABASE_URI:", app.config.get('SQLALCHEMY_DATABASE_URI'))
```

```

db = SQLAlchemy(app)

#####

# Database models (reflect your schema exactly)

#####

class Company(db.Model):

    __tablename__ = 'company'

    id = db.Column(db.Integer, primary_key=True)

    name = db.Column(db.String(255), nullable=False)

    created_at = db.Column(db.DateTime, nullable=False, server_default=func.now())


class Warehouse(db.Model):

    __tablename__ = 'warehouse'

    id = db.Column(db.Integer, primary_key=True)

    company_id = db.Column(db.Integer, db.ForeignKey('company.id', ondelete='CASCADE'),
nullable=False)

    name = db.Column(db.String(255), nullable=False)

    location = db.Column(db.String(255))

    created_at = db.Column(db.DateTime, nullable=False, server_default=func.now())


class Product(db.Model):

    __tablename__ = 'product'

    id = db.Column(db.Integer, primary_key=True)

    company_id = db.Column(db.Integer, db.ForeignKey('company.id', ondelete='CASCADE'),
nullable=False)

    name = db.Column(db.String(255), nullable=False)

    sku = db.Column(db.String(64), nullable=False)

    price = db.Column(db.Numeric(12, 2), nullable=False)

    is_bundle = db.Column(db.Boolean, nullable=False, default=False) # use this for threshold join


class Supplier(db.Model):

    __tablename__ = 'supplier'

```

```
id = db.Column(db.Integer, primary_key=True)

name = db.Column(db.String(255), nullable=False)

contact_info = db.Column(db.String(255)) # used as contact_email in response

created_at = db.Column(db.DateTime, nullable=False, server_default=func.now())
```

```
class ProductSupplier(db.Model):
```

```
    __tablename__ = 'product_supplier'

    product_id = db.Column(db.Integer, db.ForeignKey('product.id', ondelete='CASCADE'),
primary_key=True)

    supplier_id = db.Column(db.Integer, db.ForeignKey('supplier.id', ondelete='CASCADE'),
primary_key=True)
```

```
class Inventory(db.Model):
```

```
    __tablename__ = 'inventory'

    id = db.Column(db.Integer, primary_key=True)

    product_id = db.Column(db.Integer, db.ForeignKey('product.id', ondelete='CASCADE'),
nullable=False)

    warehouse_id = db.Column(db.Integer, db.ForeignKey('warehouse.id', ondelete='CASCADE'),
nullable=False)

    quantity = db.Column(db.Integer, nullable=False, default=0)
```

```
class Sale(db.Model):
```

```
    __tablename__ = 'sales'

    id = db.Column(db.Integer, primary_key=True)

    product_id = db.Column(db.Integer, db.ForeignKey('product.id'), nullable=False)

    warehouse_id = db.Column(db.Integer, db.ForeignKey('warehouse.id'), nullable=False)

    quantity = db.Column(db.Integer, nullable=False)

    sale_date = db.Column(db.DateTime, nullable=False)
```

```
class ProductTypeThreshold(db.Model):
```

```
    __tablename__ = 'product_type_threshold'

    product_type = db.Column(Boolean, primary_key=True) # BOOLEAN: True for bundles, False for
others
```

```
threshold = db.Column(db.Integer, nullable=False)
```

```
#####
```

```
# Low-stock alerts endpoint
```

```
#####
```

```
@app.route('/api/companies/<int:company_id>/alerts/low-stock', methods=['GET'])
```

```
def get_low_stock_alerts(company_id):
```

```
    """
```

```
    Get low-stock alerts for a company.
```

```
    ---
```

```
    tags:
```

```
        - Alerts
```

```
    parameters:
```

```
        - name: company_id
```

```
            in: path
```

```
            type: integer
```

```
            required: true
```

```
            description: The ID of the company to get low stock alerts for
```

```
    responses:
```

```
        200:
```

```
            description: List of low-stock alerts with supplier information
```

```
            schema:
```

```
                type: object
```

```
                properties:
```

```
                    alerts:
```

```
                        type: array
```

```
                        items:
```

```
                            type: object
```

```
                            properties:
```

product_id:

type: integer

example: 123

product_name:

type: string

example: "Widget A"

sku:

type: string

example: "WID-001"

warehouse_id:

type: integer

example: 456

warehouse_name:

type: string

example: "Main Warehouse"

current_stock:

type: integer

example: 5

threshold:

type: integer

example: 20

days_until_stockout:

type: integer

example: 12

nullable: true

supplier:

type: object

nullable: true

properties:

id:

type: integer

```

        example: 789
    name:
        type: string
        example: "Supplier Corp"
    contact_email:
        type: string
        example: "orders@supplier.com"
    total_alerts:
        type: integer
        example: 1
404:
    description: Company not found
    schema:
        type: object
        properties:
            error:
                type: string
                example: "Company not found"
"""

# Verify company exists
company = Company.query.get(company_id)
if not company:
    return jsonify({"error": "Company not found"}), 404

# Define recent sales window (last 30 days)
recent_period_start = datetime.utcnow() - timedelta(days=30)

# 1. Products with recent sales in company
recent_sales_subq = (

```

```

db.session.query(Sale.product_id)
.join(Product, Product.id == Sale.product_id)
.filter(
    Product.company_id == company_id,
    Sale.sale_date >= recent_period_start
)
.distinct()
.subquery()
)

```

2. Inventories with product info, warehouse info, and thresholds

```

inventories = (
    db.session.query(
        Inventory.id.label("inventory_id"),
        Product.id.label("product_id"),
        Product.name.label("product_name"),
        Product.sku,
        Product.is_bundle,
        Warehouse.id.label("warehouse_id"),
        Warehouse.name.label("warehouse_name"),
        Inventory.quantity.label("current_stock"),
        ProductTypeThreshold.threshold
    )
    .join(Product, Inventory.product_id == Product.id)
    .join(Warehouse, Inventory.warehouse_id == Warehouse.id)
    .join(ProductTypeThreshold, Product.is_bundle == ProductTypeThreshold.product_type) #
Boolean join here
    .filter(
        Product.company_id == company_id,
        Warehouse.company_id == company_id,
        Inventory.quantity <= ProductTypeThreshold.threshold,

```

```
        Inventory.product_id.in_(recent_sales_subq) # Only products with recent sales
    )
    .all()
)
```

3. Calculate average daily sales per product (last 30 days)

```
sales_agg = (
    db.session.query(
        Sale.product_id,
        func.sum(Sale.quantity).label("total_quantity")
    )
    .join(Product, Product.id == Sale.product_id)
    .filter(
        Product.company_id == company_id,
        Sale.sale_date >= recent_period_start
    )
    .group_by(Sale.product_id)
    .all()
)

avg_daily_sales_map = {s.product_id: s.total_quantity / 30 for s in sales_agg}
```

4. Prepare alert list

```
alerts = []
```

for inv in inventories:

```
    avg_daily_sales = avg_daily_sales_map.get(inv.product_id, 0)
```

```
    if avg_daily_sales > 0:
```

```
        days_until_stockout = int(inv.current_stock / avg_daily_sales)
```

```
    else:
```

```
        days_until_stockout = None
```

Fetch first linked supplier info (optional)


```
supplier = (  
    db.session.query(Supplier)  
    .join(ProductSupplier, Supplier.id == ProductSupplier.supplier_id)  
    .filter(ProductSupplier.product_id == inv.product_id)  
    .first()  
)  
  
supplier_info = None  
if supplier:  
    supplier_info = {  
        "id": supplier.id,  
        "name": supplier.name,  
        "contact_email": supplier.contact_info or ""  
    }  
  
alert = {  
    "product_id": inv.product_id,  
    "product_name": inv.product_name,  
    "sku": inv.sku,  
    "warehouse_id": inv.warehouse_id,  
    "warehouse_name": inv.warehouse_name,  
    "current_stock": inv.current_stock,  
    "threshold": inv.threshold,  
    "days_until_stockout": days_until_stockout,  
    "supplier": supplier_info  
}  
alerts.append(alert)  
  
return jsonify({  
    "alerts": alerts,  
    "total_alerts": len(alerts)
```

```
}), 200
```

```
#####
```

```
# Main entrypoint
```

```
#####
```

```
if __name__ == '__main__':
```

```
    # For production, use a WSGI server like gunicorn instead
```

```
    app.run(debug=True)
```

Explanation of Approach

1. **Validate Company:** Ensure the requested company exists, otherwise return 404.
2. **Recent Sales Filter:**
Query recent sales (last 30 days) to get products with sales activity. This filters out products without recent demand to avoid unnecessary alerts.
3. **Inventory & Threshold Filtering:**
Join inventory with products, warehouses, and product-type thresholds, filter for inventory levels below or equal to thresholds.
4. **Average Daily Sales Calculation:**
Calculate average daily sales per product over 30 days. This helps estimate how many days before stock runs out.
5. **Supplier Information:**
Fetch supplier details linked to each product for reordering contact info.
6. **Return Structured JSON Alert List:**
Matches the expected response format with detailed info.

Edge Cases Considered

- Company id invalid → returns 404.
- Products with no recent sales → excluded.
- Inventory quantity exactly equal to threshold → included.
- Products without any supplier → supplier info is null.
- No inventory or no low stock products → returns empty alerts list.
- Division by zero avoided by checking average daily sales > 0.
- Multi-warehouse handled by joining inventories per warehouse.

General Reasoning & Approach

- **Problem-solving:** Systematically identified input, transactional, and business logic issues; prioritized data integrity and maintainability.
- **Database:** Emphasized scalable, normalized design with attention to multi-tenancy and auditability.
- **API:** Focused on robust filtering, clear business rules, and comprehensive error handling.
- **Communication:** Asked clarifying questions to expose requirement gaps, demonstrating anticipation of real-world complexities.
- **Maintaining best practices:** Transactions, validation, appropriate HTTP status codes, modular design.